A Step in the Right Direction: Combining Point Forecasts with Foundation-Model Signals

Anonymous Author(s)
Affiliation
Address
email

Abstract

Point forecast accuracy typically deteriorates as the forecast horizon increases, but considerably less is known about how directional accuracy varies across horizons. In many operational and financial settings, the direction of change is itself decision relevant. Recent time series foundation models offer new opportunities to study directional prediction. We introduce a Directional Prediction Framework. It reformulates forecasting datasets as directional prediction tasks. Using directional accuracy as a common evaluation criterion, the framework enables direct comparison between point forecasting methods and time series classifiers. On the M4 Daily subset, Mantis exhibits higher directional accuracy at some longer forecast horizons. We further present a post hoc proof of concept analysis in which directional adjustment is associated with lower point forecast error.

1 Introduction

Point forecast accuracy typically deteriorates as the forecast horizon increases (Hyndman and Athanasopoulos, 2021). However, relatively little is known about how directional accuracy varies across horizons. This gap raises a natural question: *does directional accuracy exhibit the same horizon dependent behaviour as point forecast accuracy?*

In many applications, decisions depend less on the exact future value of a quantity than on whether it will increase or not increase. Although directional accuracy has long been recognised as an important dimension of forecast evaluation (Leitch and Tanner, 1991; Granger and Pesaran, 2000), direction is typically evaluated as a by product of point forecasts rather than modelled as a prediction target. Forecasting datasets are not structured for directional prediction, complicating comparisons between point forecasts and classifier predictions. Forecast combinations can improve accuracy (Bates and Granger, 1969), but whether signals from different model families can do the same remains unclear.

To investigate this possibility, we propose a Directional Prediction Framework that reformulates forecasting datasets into time series classification tasks with directional labels. Using directional accuracy as a common evaluation criterion, the framework enables direct comparison between point forecasting methods and time series classifiers, including foundation model based approaches. With the M4 dataset, we descriptively investigate how directional accuracy evolves across forecast horizons. We then conduct a post hoc proof of concept analysis of whether learned directional signals can improve point forecast accuracy through directional adjustment.

Our results suggest that directional accuracy does not necessarily follow the same horizon dependent pattern as point forecast accuracy. On the M4 Daily subset, Mantis exhibits higher directional accuracy at some longer forecast horizons. Among the evaluated frequencies, Daily exhibits the largest observed directional adjustment improvement, although its magnitude depends on the forecasting setting. Together, these findings suggest that directional prediction merits investigation as a task in its own right.

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

Key contributions include: **(i)** We introduce a Directional Prediction Framework that reformulates forecasting datasets as time series classification tasks with directional labels. Directional accuracy provides a common criterion for directly comparing point forecasting methods and time series classifiers, including foundation model based approaches; **(ii)** We document descriptive horizon dependent differences in directional accuracy. On the M4 Daily subset, Mantis exhibits higher directional accuracy at some longer forecast horizons; and **(iii)** We provide post hoc proof of concept evidence that learned directional signals can complement point forecasts through directional adjustment.

2 Related Work

Directional accuracy measures whether a prediction correctly identifies the direction of change of a target variable. It has long been recognised as an important dimension of forecast evaluation because many decisions depend more on anticipating future movements than on minimising point forecast error (Leitch and Tanner, 1991; Granger and Pesaran, 2000). Despite its practical importance, directional prediction has traditionally been treated as a secondary outcome of point forecasting rather than as a task in its own right (De Gooijer and Hyndman, 2006; Hyndman and Koehler, 2006). Recent studies underscore the need for precise directional evaluation protocols, as ambiguous formulations can distort directional accuracy estimates (McCarthy and Snudden, 2025).

Recent advances in **time series classification** (TSC) have produced effective methods for extracting or learning discriminative temporal representations, including ROCKET and InceptionTime, together with standardised evaluation frameworks based on benchmark repositories such as the UCR Archive (Bagnall et al., 2017; Dau et al., 2019; Dempster et al., 2020; Ismail Fawaz et al., 2020). At the same time, **time series foundation models** have created new opportunities for forecasting and classification. Chronos-2 uses group attention for zero-shot univariate, multivariate, and covariate informed forecasting (Ansari et al., 2025). Mantis is a lightweight transformer based foundation model for time series classification trained with self supervised contrastive learning (Feofanov et al., 2025).

These developments motivate a unified view of directional prediction. Although point forecasting methods and time series classifiers, including foundation model based approaches, all capture temporal information, they are typically evaluated using different criteria and therefore remain difficult to compare directly. Our framework enables their comparison using directional accuracy as a common evaluation criterion across forecast horizons.

3 Method

We introduce a Directional Prediction Framework that reformulates each forecasting task as a binary classification problem: predicting whether a future observation increases relative to the forecast origin. The framework evaluates point forecasting methods and time series classifiers using directional accuracy. When the classifier and forecast imply different labels, the classifier label determines whether the upward or downward adjustment factor is applied to the point forecast.

Forecasting notation. Let $\mathbf{y} = (y_1, y_2, \ldots, y_T)$ be the observed time series, with forecast origin T and final observed value y_T. At horizon h, y_{T+h} is the future observation and $\hat{y}_{T+h|T}$ is its point forecast. Let H denote the maximum forecast horizon, with $h \in \{1, \ldots, H\}$.

Binary sign convention. Following McCarthy and Snudden (2025), define $\mathrm{sign}(y) = 1$ if $y > 0$, and $\mathrm{sign}(y) = -1$ if $y \leq 0$.

Directional accuracy. At forecast origin T and horizon h, directional accuracy is defined as

$$d_{T,h} = \mathbb{I}\left[\mathrm{sign}(y_{T+h} - y_T) = \mathrm{sign}(\hat{y}_{T+h|T} - y_T)\right]. \tag{1}$$

Both changes are evaluated relative to the forecast origin y_T. For a collection of N time series, let $i \in \{1, \ldots, N\}$ index the individual series. Across series directional accuracy is $\mathrm{DA}_{T,h} = \frac{1}{N}\sum_{i=1}^{N} d_{i,T,h}$, the mean of the series level indicators (Appendix C).

Directional class labels. At forecast origin T and horizon h, define

$$z_{T,h} = \mathbb{I}\left[\mathrm{sign}(y_{T+h} - y_T) = 1\right], \tag{2}$$

where $z_{T,h} = 1$ denotes an increase from y_T to y_{T+h}, and $z_{T,h} = 0$ denotes a non increase.

Directional Prediction Framework. The framework proceeds in four steps (Figure 1).

1. **Classifier predicted directional class label.** At forecast origin T and horizon h, a time series classifier produces $\hat{z}_{T,h} \in \{0, 1\}$, where $\hat{z}_{T,h} = 1$ denotes a predicted increase relative to y_T, and $\hat{z}_{T,h} = 0$ denotes a predicted non increase.
2. **Forecast implied directional class label.** The point forecast implies
$$\hat{z}^{\mathrm{F}}_{T,h} = \mathbb{I}\left[\mathrm{sign}(\hat{y}_{T+h|T} - y_T) = 1\right]. \tag{3}$$
3. **Directional adjustment.** The adjustment factor is
$$\gamma_{T,h} = \begin{cases} 1, & \hat{z}_{T,h} = \hat{z}^{\mathrm{F}}_{T,h}, \\ \lambda_{\mathrm{Up}}, & \hat{z}_{T,h} \neq \hat{z}^{\mathrm{F}}_{T,h} \ \wedge \ \hat{z}_{T,h} = 1, \\ \lambda_{\mathrm{Down}}, & \hat{z}_{T,h} \neq \hat{z}^{\mathrm{F}}_{T,h} \ \wedge \ \hat{z}_{T,h} = 0. \end{cases} \tag{4}$$
where $\lambda_{\mathrm{Up}} \geq 1$ and $0 < \lambda_{\mathrm{Down}} \leq 1$ are multiplicative adjustment parameters.
4. **Adjusted forecast.** The adjusted point forecast is
$$\tilde{y}_{T+h|T} = \gamma_{T,h}\hat{y}_{T+h|T}. \tag{5}$$

The adjustment in equation (5) is used as a post hoc exploratory proof of concept to assess whether learned directional information can improve point forecast accuracy, rather than as an optimal adjustment rule. We use a multiplicative adjustment because M4 target values are positive.

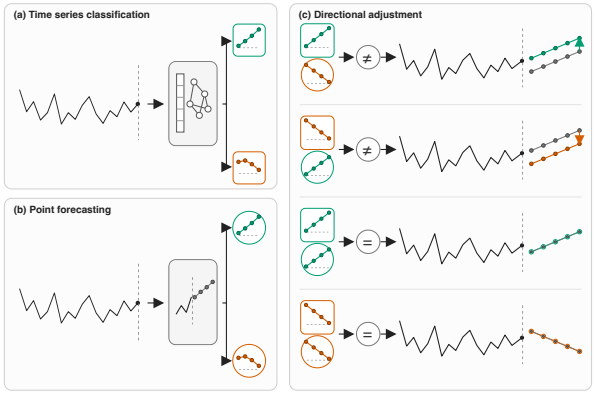


Figure 1: Directional Prediction Framework. Squares and circles denote classifier-predicted and forecast implied labels, respectively; green and orange indicate increase and non increase. Agreement preserves the forecast; disagreement applies the classifier indicated adjustment.

4 Experiments

We evaluate the Directional Prediction Framework on the M4 dataset (Makridakis et al., 2018), comprising 100,000 time series across six frequencies and multiple forecast horizons. Following Section 3, we reformulate the forecasting tasks as directional prediction tasks. We evaluate point forecasting methods and time series classifiers, including foundation model based approaches, using directional accuracy. We compare the M4 winning SMYL (Smyl, 2020) and second ranked FFORMA (Montero-Manso et al., 2020) point forecasts with Chronos-2 (Ansari et al., 2025). The classifier comparison includes 1-NN DTW, 1-NN ED, Rotation Forest, XGBoost, ROCKET (Dempster et al., 2020), InceptionTime (Ismail Fawaz et al., 2020), and Mantis (Feofanov et al., 2025). The conventional classifiers span distance based, feature based, ensemble, kernel based, and deep learning approaches. Chronos-2 uses released weights without fine tuning but has documented M4 pretraining overlap. Mantis uses a frozen encoder with a 200 tree Random Forest. Implementation and dataset construction details are provided in Appendices C–D. Directional prediction is evaluated independently at each horizon. Tables 1 and 2 report directional accuracy at the terminal official horizon $h = H$, whereas sMAPE, MASE, and OWA are computed over the complete official forecast horizon. Forecast implied directional labels are calculated using equation (3). We assess whether learned directional signals improve point forecast accuracy through a post hoc parameter grid analysis on the official M4 evaluation set (Appendix B). Accordingly, adjusted results are interpreted as exploratory proof of concept estimates rather than held out performance estimates.

5 Results

Figure 2 descriptively shows how directional accuracy varies across horizons on the M4 Daily subset, with the highest observed values occurring at some longer horizons.

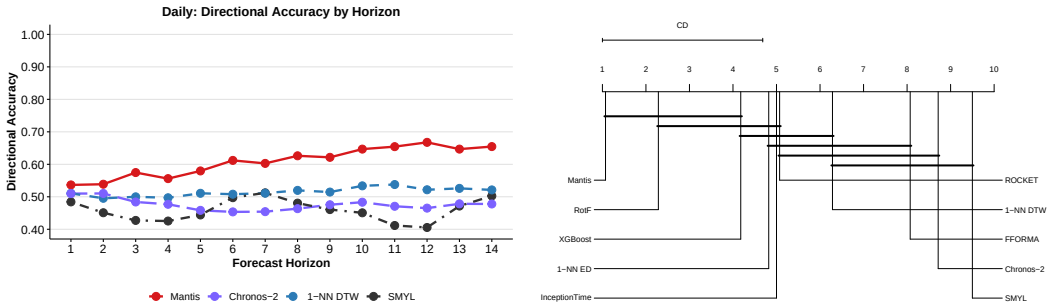


Figure 2: Directional prediction results for the M4 Daily subset. **Left:** Directional accuracy by forecast horizon. **Right:** Descriptive mean ranks across the 14 Daily forecast horizons, with lower ranks indicating better performance.

Mantis achieves the best average rank across the Daily forecast horizons. This result is not driven solely by an individual horizon. Across the remaining frequencies, directional profiles are heterogeneous, and no single method dominates (Appendix A). Additional details on the rank diagram and reproducibility are provided in Appendix D.

Table 1 reports the forecasting value of directional adjustment on the Daily subset. Relative to their corresponding unadjusted forecasts, the Daily OWA reductions are 4.01% for SMYL–Mantis and 5.52% for Chronos-2–Mantis. The 7.64% value in the table instead compares Chronos-2–Mantis with SMYL. The infeasible SMYL–Oracle achieves a substantially larger 47.34% OWA improvement relative to SMYL, indicating considerable ex post headroom within the scaling construction rather than the performance of a feasible forecasting method.

Table 1: Daily forecasting performance. DA is reported at $h = 14$; sMAPE, MASE, and OWA are reported over all official horizons. Bold indicates the best feasible value in each column.

Method	**sMAPE↓**	**MASE↓**	**DA↑**	**OWA↓**	**Improvement over SMYL (%)**		
					sMAPE↑	**MASE↑**	**OWA↑**
Chronos-2	3.09	3.38	0.527	1.023	2.61	1.88	2.24
FFORMA	3.10	3.34	0.510	1.019	2.29	2.95	2.62
SMYL	3.17	3.45	0.502	1.046	0.00	0.00	0.00
SMYL–Mantis	3.05	3.30	**0.670**	1.004	3.69	4.33	4.01
Chronos-2–Mantis	**2.95**	**3.16**	0.665	**0.966**	**6.86**	**8.41**	**7.64**
SMYL–Oracle	1.72	1.76	0.859	0.551	45.64	49.02	47.34

6 Conclusion

We introduced a Directional Prediction Framework that enables point forecasting methods and time series classifiers to be evaluated using directional accuracy and uses learned directional signals to adjust point forecasts. On the M4 Daily subset, Mantis exhibits higher directional accuracy at some longer horizons, while directional accuracy profiles vary across frequencies. Mantis based directional adjustment also lowers the OWA of SMYL and Chronos-2 on the Daily subset. Gains vary across frequencies and forecasting settings. These findings provide exploratory proof of concept evidence that learned directional representations can complement point forecasts. Time series foundation models therefore offer opportunities to study directional prediction as a complementary forecasting task. The evidence is limited to the M4 dataset and the methods considered. Future work should evaluate additional datasets and representations, investigate mechanisms underlying directional predictability, and develop more adaptive adjustment strategies. All comparisons are descriptive and conditional on the reported experimental design (Appendix D).

References

Abdul Fatir Ansari, Oleksandr Shchur, Jaris Küken, et al. Chronos-2: From univariate to universal forecasting. *arXiv preprint arXiv:2510.15821*, 2025. doi: 10.48550/arXiv.2510.15821. URL `https://arxiv.org/abs/2510.15821`.

Anthony Bagnall, Jason Lines, Aaron Bostrom, James Large, and Eamonn Keogh. The great time series classification bake off: a review and experimental evaluation of recent algorithmic advances. *Data mining and knowledge discovery*, 31(3):606–660, 2017. ISSN 1384-5810.

J. M. Bates and C. W. J. Granger. The combination of forecasts. *Operational Research Quarterly*, 20 (4):451–468, 1969. doi: 10.1057/jors.1969.103.

Borja Calvo and Guzmán Santafé. `scmamp`: Statistical comparison of multiple algorithms in multiple problems. *The R Journal*, 8(1):248–256, 2016. doi: 10.32614/RJ-2016-017.

Hoang Anh Dau, Anthony Bagnall, Kaveh Kamgar, Chin-Chia Michael Yeh, Yan Zhu, Shaghayegh Gharghabi, Chotirat Ann Ratanamahatana, and Eamonn Keogh. The ucr time series archive, 2019. URL `https://arxiv.org/abs/1810.07758`.

Jan G. De Gooijer and Rob J. Hyndman. 25 years of time series forecasting. *International journal of forecasting*, 22(3):443–473, 2006. ISSN 0169-2070.

Angus Dempster, François Petitjean, and Geoffrey I. Webb. Rocket: exceptionally fast and accurate time series classification using random convolutional kernels. *Data mining and knowledge discovery*, 34(5):1454–1495, 2020. ISSN 1384-5810.

Vasilii Feofanov, Songkang Wen, Marius Alonso, Romain Ilbert, Hongbo Guo, Malik Tiomoko, Lujia Pan, Jianfeng Zhang, and Ievgen Redko. Mantis: Lightweight calibrated foundation model for user-friendly time series classification. *arXiv preprint arXiv:2502.15637*, 2025.

Clive W.J. Granger and M. Hashem Pesaran. Economic and statistical measures of forecast accuracy. *Journal of forecasting*, 19(7):537–560, 2000. ISSN 0277-6693.

Rob J. Hyndman and George Athanasopoulos. *Forecasting: Principles and Practice*. OTexts, Melbourne, Australia, 3rd edition, 2021. ISBN 9780987507136.

Rob J. Hyndman and Anne B. Koehler. Another look at measures of forecast accuracy. *International journal of forecasting*, 22(4):679–688, 2006. ISSN 0169-2070.

Hassan Ismail Fawaz, Benjamin Lucas, Germain Forestier, Charlotte Pelletier, Daniel F. Schmidt, Jonathan Weber, Geoffrey I. Webb, Lhassane Idoumghar, Pierre-Alain Muller, and François Petitjean. Inceptiontime: Finding alexnet for time series classification. *Data mining and knowledge discovery*, 34(6):1936–1962, 2020. ISSN 1384-5810.

Gordon Leitch and J. Ernest Tanner. Economic forecast evaluation: Profits versus the conventional error measures. *The American economic review*, 81(3):580–590, 1991. ISSN 0002-8282.

Markus Löning, Anthony Bagnall, Sajaysurya Ganesh, Viktor Kazakov, Jason Lines, and Franz J Király. sktime: A unified interface for machine learning with time series. *arXiv.org*, 2019. ISSN 2331-8422.

Spyros Makridakis, Evangelos Spiliotis, and Vassilios Assimakopoulos. The m4 competition: Results, findings, conclusion and way forward. *International journal of forecasting*, 34(4):802–808, 2018. ISSN 0169-2070.

Martin McCarthy and Stephen Snudden. Predictable by construction: assessing forecast directional accuracy of temporal aggregates. *Applied economics*, pages 1–16, 2025. ISSN 0003-6846.

Pablo Montero-Manso, Carla Netto, and Thiyanga Talagala. *M4comp2018: Data from the M4-Competition*, 2018. URL `https://github.com/carlanetto/M4comp2018`. R package version 0.2.0.

Pablo Montero-Manso, George Athanasopoulos, Rob J. Hyndman, and Thiyanga S. Talagala. Fforma: Feature-based forecast model averaging. *International journal of forecasting*, 36(1):86–92, 2020. ISSN 0169-2070.

Slawek Smyl. A hybrid method of exponential smoothing and recurrent neural networks for time series forecasting. *International journal of forecasting*, 36(1):75–85, 2020. ISSN 0169-2070.

A Forecasting Value of Directional Information

This appendix examines whether learned directional signals provide additional value for point forecasting across the six M4 frequencies. Figure 3 illustrates two directional adjustments when classifier predicted and forecast implied directions disagree. Table 2 summarises terminal horizon directional accuracy and full horizon point forecast accuracy across frequencies.

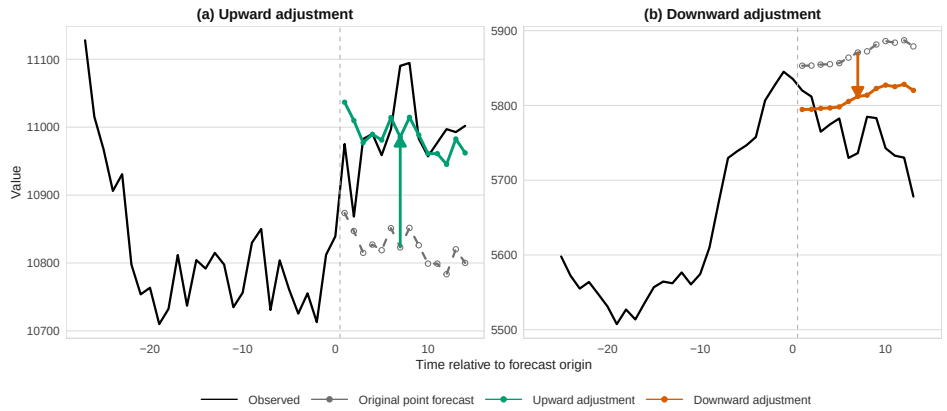


Figure 3: Directional adjustment examples from M4. Mantis adjusts the SMYL forecast upward for Daily series D1602 (left) and the Chronos-2 forecast downward for Weekly series W39 (right).

Table 2: Directional accuracy at the terminal official horizon and point forecast accuracy over the complete official forecast horizon, by M4 frequency.

Panel A. Directional accuracy at $h = H$ (DA) ↑

Method	Hourly	Daily	Weekly	Monthly	Quarterly	Yearly	All
Chronos-2	0.802	0.527	0.705	0.664	0.657	0.676	0.660
FFORMA	0.688	0.510	0.649	0.694	0.706	0.765	0.705
SMYL	0.761	0.502	0.635	**0.722**	0.717	0.784	0.726
SMYL–Mantis	0.766	**0.670**	0.652	**0.722**	**0.721**	**0.786**	**0.734**
Chronos-2–Mantis	**0.809**	0.665	**0.708**	0.665	0.693	0.725	0.686
Mantis	0.652	0.658	0.593	0.521	0.699	0.778	0.629
SMYL–Oracle	0.829	0.859	0.841	0.859	0.861	0.909	0.871

Panel B. Overall Weighted Average (OWA) ↓

Method	Hourly	Daily	Weekly	Monthly	Quarterly	Yearly	All
Chronos-2	0.393	1.023	0.712	0.888	0.911	0.891	0.895
FFORMA	0.484	1.019	0.796	0.858	0.847	0.799	0.838
SMYL	0.440	1.046	0.851	**0.836**	0.847	0.778	0.821
SMYL–Mantis	0.438	1.004	0.839	**0.836**	**0.847**	**0.778**	**0.819**
Chronos-2–Mantis	**0.390**	**0.966**	**0.695**	0.888	0.898	0.864	0.880
SMYL–Oracle	0.356	0.551	0.540	0.557	0.502	0.383	0.479

Panel C. OWA improvement relative to SMYL (%) ↑

Method	Hourly	Daily	Weekly	Monthly	Quarterly	Yearly	All
Chronos-2	10.78	2.24	16.29	-6.29	-7.57	-14.51	-9.01
FFORMA	-9.92	2.62	6.49	-2.65	0.03	-2.69	-2.05
SMYL	0.00	0.00	0.00	**0.00**	0.00	0.00	0.00
SMYL–Mantis	0.57	4.01	1.42	**0.00**	**0.07**	**0.02**	**0.25**
Chronos-2–Mantis	**11.31**	**7.64**	**18.27**	-6.23	-6.00	-11.02	-7.19
SMYL–Oracle	19.03	47.34	36.52	33.40	40.69	50.76	41.61

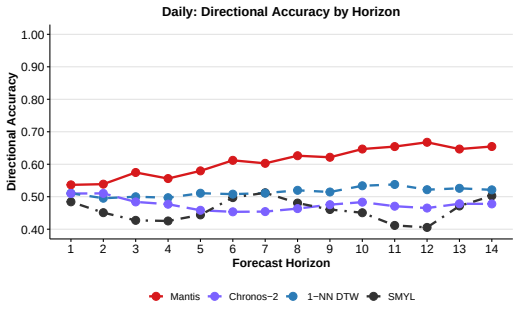
Notes.

(i) Bold indicates the best feasible value in each column; Panel C percentages use unrounded OWA values.
(ii) Terminal H for Hourly through Yearly is 48, 14, 13, 18, 8, and 6, respectively.
(iii) All DA is weighted by series count; All OWA is recomputed from weighted sMAPE and MASE relative to Naive2.
(iv) Adjusted results use parameters selected on the evaluation set and are exploratory; SMYL Oracle uses realised outcomes and is infeasible.

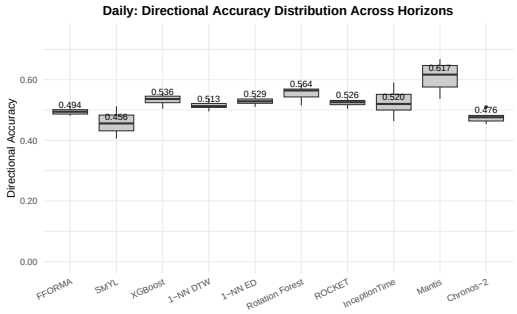
The Daily subset offers the clearest demonstration. Mantis based adjustment increases Chronos-2 terminal horizon directional accuracy from 0.527 to 0.665 and reduces OWA from 1.023 to 0.966. The same adjustment increases SMYL terminal horizon directional accuracy from 0.502 to 0.670 and reduces OWA from 1.046 to 1.004. These adjusted OWA results use the post hoc parameter grid analysis described in Appendix B and are therefore exploratory proof of concept estimates rather than held out performance estimates.

The Hourly and Weekly results extend this pattern beyond Daily data: Chronos-2–Mantis achieves the lowest non oracle OWA at all three frequencies. Comparable gains do not occur uniformly for the lower frequency subsets or in the aggregate results, indicating that the value of learned directional information is setting dependent.

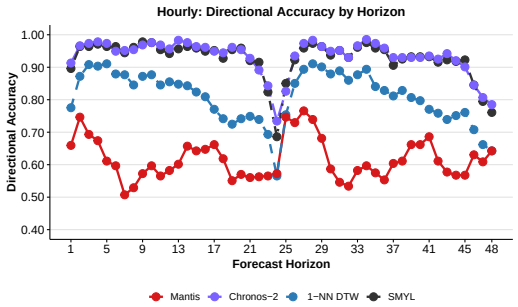
Figures 2, 4, and 5 show the corresponding horizon structure. On the Daily subset, Mantis directional accuracy increases from 0.537 at the first horizon to 0.668 at horizon 12; the remaining frequencies exhibit more setting specific profiles. These descriptive horizon profiles vary across frequencies and need not follow the corresponding patterns of point forecast accuracy.



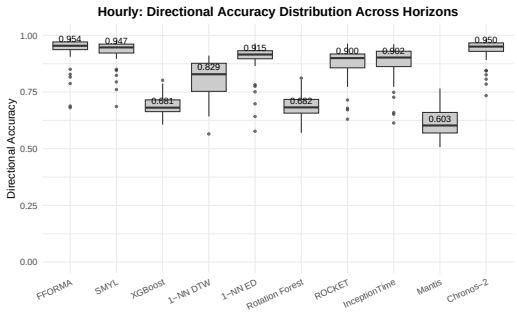
(a) Daily: horizon profile



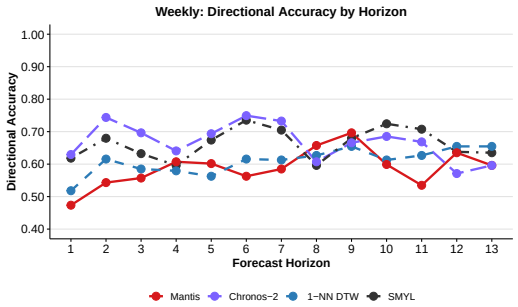
(b) Daily: distribution



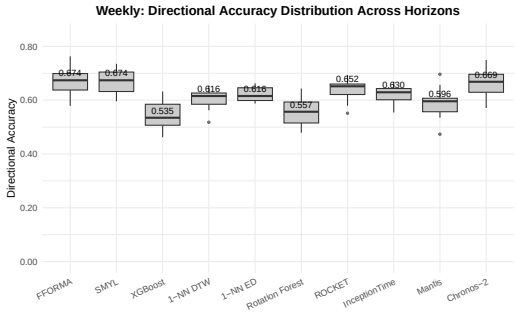
(c) Hourly: horizon profile



(d) Hourly: distribution

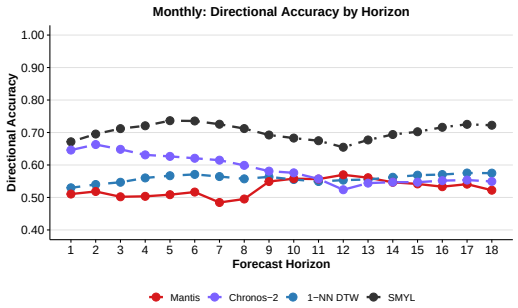


(e) Weekly: horizon profile

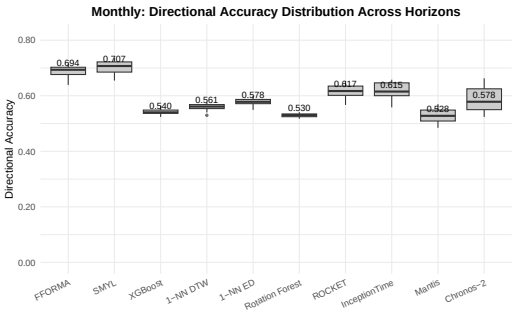


(f) Weekly: distribution

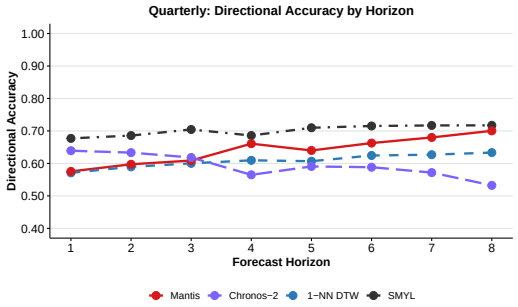
Figure 4: Directional accuracy across forecast horizons for the M4 Daily, Hourly, and Weekly subsets. **Left:** horizon profiles for selected methods. **Right:** box plots for all methods; labels report medians.



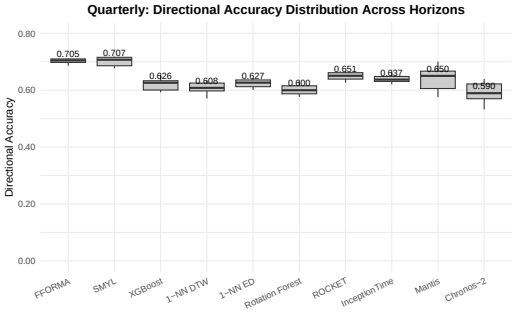
(a) Monthly: horizon profile



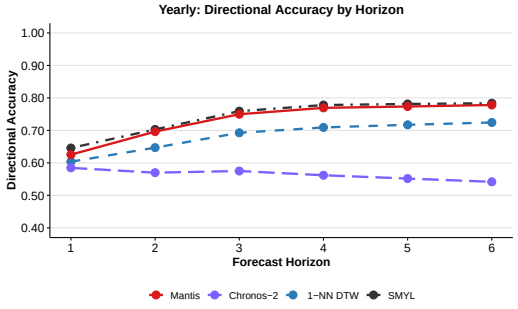
(b) Monthly: distribution



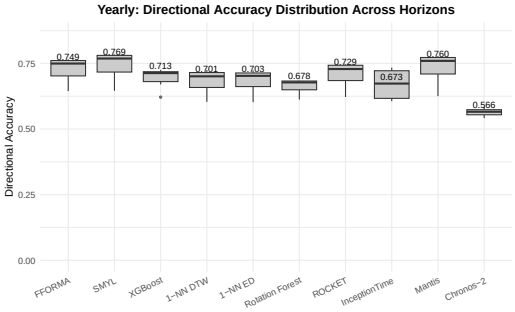
(c) Quarterly: horizon profile



(d) Quarterly: distribution



(e) Yearly: horizon profile



(f) Yearly: distribution

Figure 5: Directional accuracy across forecast horizons for the M4 Monthly, Quarterly, and Yearly subsets. **Left:** horizon profiles for selected methods. **Right:** box plots for all methods; labels report medians.

SMYL–Oracle is the infeasible ex post scaling diagnostic. For each series, it selects the multiplier $m \in \{0.500, 0.501, \ldots, 1.500\}$ that minimises SMYL sMAPE over that series' realised official forecast horizon. On the Daily subset, it reduces OWA by 47.34% relative to SMYL. The gap between this infeasible result and the feasible SMYL–Mantis result indicates headroom within the scaling construction when realised outcomes are available, but it should not be attributed solely to unexploited directional information.

Together, the results provide exploratory evidence that learned directional representations can complement point forecasts in suitable settings. The Directional Prediction Framework makes this contribution measurable across point forecasting methods, time series classifiers, and future foundation model representations.

B Sensitivity Analysis of Directional Adjustment

This appendix assesses whether directional adjustment gains are robust to variation in the adjustment parameters. For each model frequency combination, we evaluated 525 global parameter pairs, with $\lambda_{\mathrm{Up}} \in [1.00, 1.12]$ and $\lambda_{\mathrm{Down}} \in [0.90, 1.00]$, both in increments of 0.005. Each pair was evaluated using OWA on the official M4 evaluation set, and the minimum OWA pair was retained. Consequently, these adjusted results are exploratory proof of concept estimates rather than held out performance estimates. Figure 6 shows the Daily response surfaces for both adjusted forecasts.

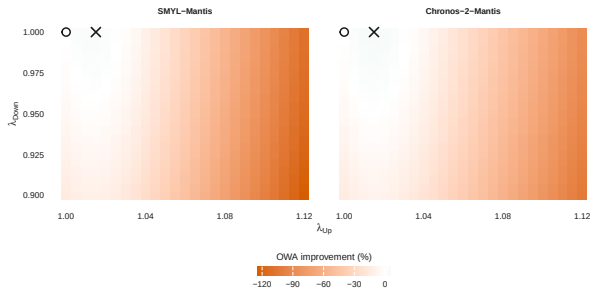


Figure 6: Directional adjustment sensitivity for the M4 Daily subset. Colour shows OWA improvement relative to the corresponding unadjusted forecast. Circles mark no adjustment, $(\lambda_{\mathrm{Up}}, \lambda_{\mathrm{Down}}) = (1, 1)$; crosses mark the minimum OWA setting in each panel.

The Daily response surfaces are smooth in the vicinity of the minimum OWA settings. Performance varies more strongly with λ_{Up} than with λ_{Down}, but improvements occur over local regions rather than at isolated parameter pairs. Table 3 summarises the minimum OWA settings and robustness measures across frequencies.

Table 3: Directional adjustment settings and robustness across M4 frequencies.

Panel A. Minimum OWA parameter settings

Frequency	Adjusted method	λ^*_{Up}	$\lambda^*_{\mathrm{Down}}$	Unadjusted OWA ↓	Minimum OWA ↓	OWA improvement (%) ↑
Hourly	Chronos-2–Mantis	1.005	1.000	0.393	**0.390**	**0.60**
Hourly	SMYL–Mantis	1.005	1.000	0.440	0.438	0.57
Daily	Chronos-2–Mantis	1.015	1.000	1.023	**0.966**	**5.52**
Daily	SMYL–Mantis	1.015	1.000	1.046	1.004	4.01
Weekly	Chronos-2–Mantis	1.015	0.990	0.712	**0.695**	**2.38**
Weekly	SMYL–Mantis	1.015	0.995	0.851	0.839	1.42
Monthly	Chronos-2–Mantis	1.005	1.000	0.888	0.888	**0.06**
Monthly	SMYL–Mantis	1.000	1.000	0.836	**0.836**	0.00
Quarterly	Chronos-2–Mantis	1.025	1.000	0.911	0.898	**1.45**
Quarterly	SMYL–Mantis	1.005	1.000	0.847	**0.847**	0.07
Yearly	Chronos-2–Mantis	1.065	0.995	0.891	0.864	**3.05**
Yearly	SMYL–Mantis	1.010	1.000	0.778	**0.778**	0.02

Panel B. Robustness over the parameter grid

Frequency	Adjusted method	λ_{Up} range	λ_{Down} range	Improving grid points (%) ↑
Hourly	Chronos-2–Mantis	1.000–1.005	1.000–1.000	0.19
Hourly	SMYL–Mantis	1.000–1.010	0.995–1.000	**0.38**
Daily	Chronos-2–Mantis	1.010–1.020	0.985–1.000	**9.33**
Daily	SMYL–Mantis	1.010–1.015	0.985–1.000	6.67
Weekly	Chronos-2–Mantis	1.005–1.025	0.980–1.000	**7.81**
Weekly	SMYL–Mantis	1.005–1.035	0.985–1.000	6.48
Monthly	Chronos-2–Mantis	1.000–1.015	0.975–1.000	**0.38**
Monthly	SMYL–Mantis	1.000–1.010	0.980–1.000	0.00
Quarterly	Chronos-2–Mantis	1.005–1.045	0.900–1.000	**35.43**
Quarterly	SMYL–Mantis	1.000–1.030	0.910–1.000	0.95
Yearly	Chronos-2–Mantis	1.030–1.110	0.905–1.000	**94.48**
Yearly	SMYL–Mantis	1.000–1.085	0.915–1.000	1.14

Notes.

(i) Bold indicates the better of the two methods within each frequency and outcome column.
(ii) OWA improvement is $100(1 - \mathrm{OWA}_{\mathrm{adjusted}}/\mathrm{OWA}_{\mathrm{unadjusted}})$ from unrounded values; improving grid points is the percentage of 525 pairs with OWA below the corresponding unadjusted value.
(iii) Parameter ranges are marginal extrema among grid points no more than 1% above minimum OWA; points inside the implied rectangle need not meet this criterion.
(iv) Results use parameters selected on the official M4 evaluation set and are exploratory, not held out estimates.

OWA improvement is reported relative to each method's corresponding unadjusted forecast. For Daily, the reductions are 4.01% for SMYL–Mantis relative to SMYL and 5.52% for Chronos-2–Mantis relative to Chronos-2. Daily shows the largest within model improvements, whereas Monthly adjustments provide negligible gains. Quarterly and Yearly Chronos-2–Mantis adjustments have the broadest improving regions but more modest best improvements. Thus, attainable gain and the extent of the parameter region over which improvement occurs are distinct properties. Figure 7 shows the corresponding response surfaces across frequencies and base forecasts.

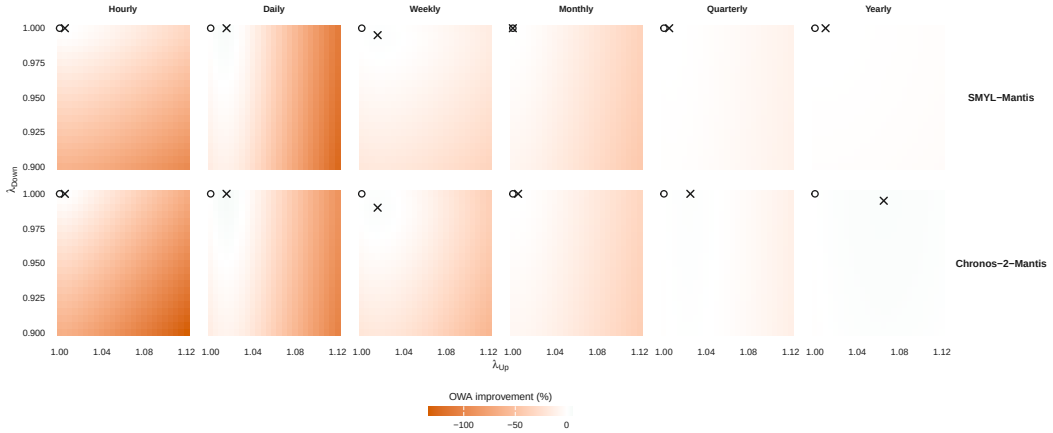


Figure 7: Directional adjustment sensitivity across M4 frequencies and base forecasts. Colour shows OWA improvement relative to the corresponding unadjusted forecast. Circles mark no adjustment, $(\lambda_{\mathrm{Up}}, \lambda_{\mathrm{Down}}) = (1, 1)$; crosses mark the minimum OWA setting in each panel.

The shape and extent of the improving regions differ across frequencies and base forecasts. Surfaces remain smooth throughout the evaluated grid. At the selected Daily setting $(\lambda_{\mathrm{Up}}, \lambda_{\mathrm{Down}}) = (1.015, 1.000)$, Mantis disagrees with SMYL and Chronos-2 on 46.94% and 38.06% of the series, respectively. Among these disagreements, the adjusted terminal forecast adopts the Mantis label in 86.90% and 76.88% of cases, respectively, while OWA decreases by 8.48% and 13.31% within the corresponding disagreement subsets. Because $\lambda_{\mathrm{Down}} = 1$, forecasts associated with downward disagreements remain unchanged.

C Directional Prediction Dataset Construction

Section 3 defines the directional labels and evaluation criterion. Forecasting datasets such as M4 are not naturally structured as directional prediction tasks. This appendix operationalises the framework by converting unequal length M4 series into fixed length, labelled instances for horizon specific binary classification. The resulting datasets follow the labelled series conventions used in the UCR Archive and TSC Bake Off studies (Dau et al., 2019; Bagnall et al., 2017). Point forecasting methods retain their native forecast generation procedures.

Directional accuracy across series. For a method applied to a dataset containing N time series, let $i \in \{1, \ldots, N\}$ denote the index of an individual series. Directional accuracy at forecast origin T and horizon h is defined as

$$\mathrm{DA}_{T,h} = \frac{1}{N} \sum_{i=1}^{N} d_{i,T,h}, \tag{6}$$

where $d_{i,T,h}$ denotes the directional accuracy indicator in equation (1) for time series i.

Construction of directional prediction tasks. Consider one rolling window of length w, ending at forecast origin T:

$$\mathbf{y}_{T,w} = (y_{T-w+1}, \ldots, y_T).$$

The procedure is repeated across rolling windows and time series.

1. **Construct rolling windows.** Following the rolling origin principle described by Hyndman and Athanasopoulos (2021), fixed length windows are generated using window length w, maximum forecast horizon H, and step size $w + H$. For the reported results, $w = (256, 64, 64, 64, 32, 16)$ for Hourly, Daily, Weekly, Monthly, Quarterly, and Yearly series, respectively.
2. **Standardise the rolling windows.** Let μ and σ be the mean and standard deviation of $\mathbf{y}_{T,w}$. If $\sigma > 0$, each observation is standardised as

$$y_t^{\mathrm{std}} = \frac{y_t - \mu}{\sigma}, \qquad t = T - w + 1, \ldots, T.$$

   Applying this transformation elementwise gives the standardised window $\mathbf{y}_{T,w}^{\mathrm{std}}$, which is the classifier input. If $\sigma = 0$, $\mathbf{y}_{T,w}^{\mathrm{std}}$ is set to the zero vector.
3. **Assign directional class labels.** For each forecast horizon $h = 1, \ldots, H$, the window is assigned the directional label $z_{T,h}$ defined in equation (2). The future observation y_{T+h} is used only to assign the label and is not included in the classifier input.
4. **Fit the classifiers.** Series yielding at least one complete length w input window followed by H outcomes were randomly assigned to training and internal hold out partitions using an 80:20 series level split. Ineligible series were omitted from classifier training without padding or imputation, and all windows from a series remained in the same partition. A separate classifier was fitted for each horizon using only windows from the training series.

This construction converts unequal length forecasting series into fixed length, scale standardised, horizon specific directional prediction datasets while preserving the information available at each forecast origin. It thereby enables time series classifiers and point forecasting methods to be evaluated on a common directional task. Figure 8 summarises the rolling window construction and resulting time series classification tasks.

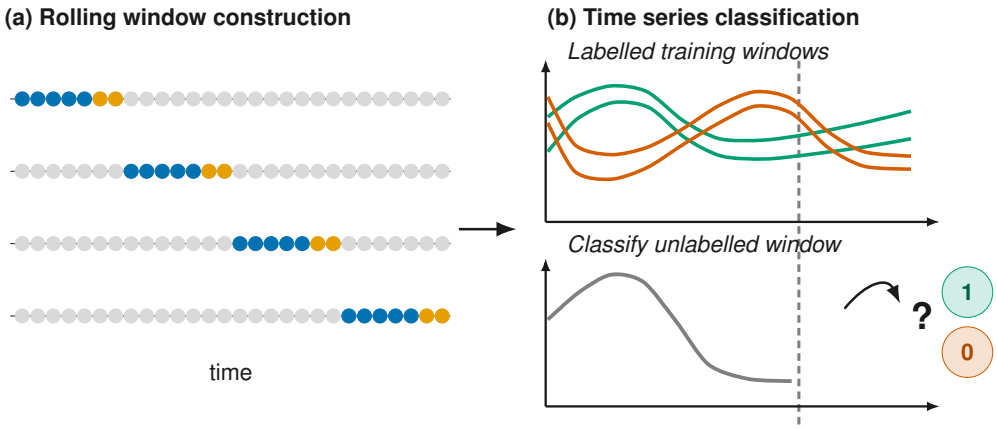


Figure 8: Construction of horizon specific directional prediction datasets. **(a)** Blue points form fixed length classifier inputs, orange points are future observations used to assign labels, and grey points lie outside each instance. **(b)** At horizon h, labelled windows train a binary classifier that predicts increase ($z_{T,h} = 1$) or non increase ($z_{T,h} = 0$) for a new window.

D Experimental Design and Reproducibility

Dataset characteristics. Experiments use the M4 dataset (Makridakis et al., 2018). Table 4 reports the number of series by frequency and domain. Figure 9 reports directional class proportions at the terminal official horizon H for each frequency, using the labels defined in equation (2).

Table 4: Number of M4 series by frequency and domain.

Frequency	Micro	Industry	Macro	Finance	Demographic	Other	Total
Yearly	6,538	3,716	3,903	6,519	1,088	1,236	23,000
Quarterly	6,020	4,637	5,315	5,305	1,858	865	24,000
Monthly	10,975	10,017	10,016	10,987	5,728	277	48,000
Weekly	112	6	41	164	24	12	359
Daily	1,476	422	127	1,559	10	633	4,227
Hourly	0	0	0	0	0	414	414
Total	25,121	18,798	19,402	24,534	8,708	3,437	100,000

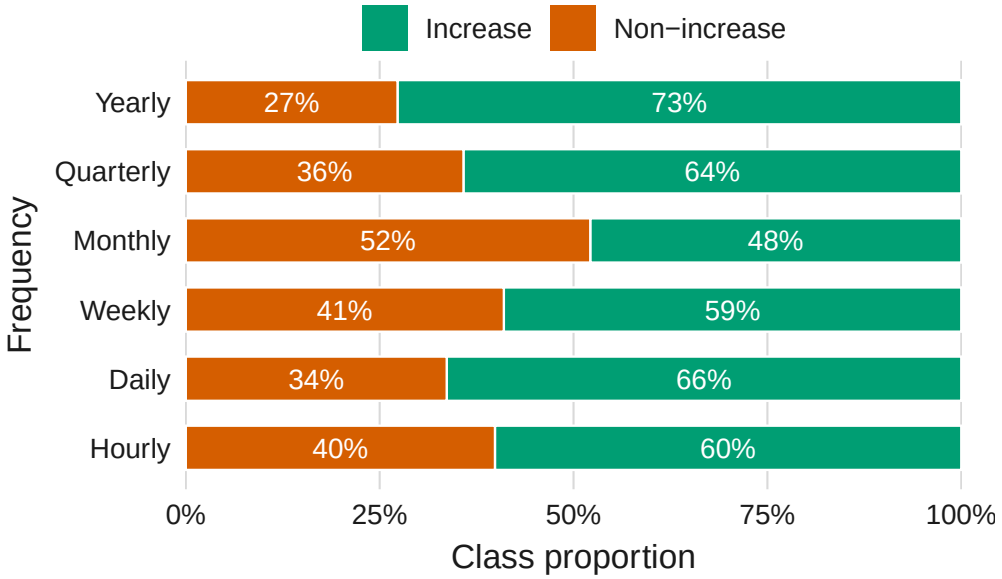


Figure 9: Directional class proportions across all series at the terminal official horizon H for each M4 frequency. Increase denotes $y_{T+H} > y_T$; non increase denotes $y_{T+H} \leq y_T$, where y_T is the final observation.

Training and evaluation protocol. Dataset construction, including the series level split, training window eligibility, and classifier fitting, is described in Appendix C. At each horizon, every M4 series contributed one directional prediction at its official final forecast origin; these all series results include both training partition and held out series, although the corresponding future observations were not used for classifier training. Directional accuracy was computed using equations (1) and (6).

Model configuration. SMYL and FFORMA forecasts were obtained from `M4comp2018` (Montero-Manso et al., 2018). Chronos-2 used the publicly available `amazon/chronos-2` checkpoint without task specific fine tuning. Its pretraining data include the M4 Daily, Hourly, Monthly, and Weekly subsets (Ansari et al., 2025); results for these frequencies therefore involve pretraining overlap. Mantis followed the authors' frozen backbone feature extraction protocol (Feofanov et al., 2025): representations were extracted without fine tuning and classified using a 200 tree Random Forest with unrestricted maximum depth. The pretraining corpus reported for Mantis does not include M4; therefore, unlike Chronos-2, Mantis has no documented direct M4 pretraining overlap. ROCKET,

InceptionTime, 1-NN DTW, 1-NN ED, and Rotation Forest used their standard `sktime` implementations (Löning et al., 2019). For each frequency, window setting, and horizon, XGBoost used a binary logistic objective; its eight hyperparameters were selected by five fold log loss Bayesian optimisation with 8 initial and 20 sequential evaluations, at most 1,500 boosting rounds, and 15 round early stopping, following the FFORMA tuning approach (Montero-Manso et al., 2020). Directional adjustment follows equations (4) and (5), using the post hoc parameter values reported in Appendix B.

Model comparison. The right panel of Figure 2 was generated using `scmamp::plotCD` with $\alpha = 0.05$ (Calvo and Santafé, 2016) and reports mean ranks across the 14 Daily forecast horizons. Because all horizons are evaluated on the same Daily series, the critical difference bar and grouping lines are not interpreted as evidence of statistical significance. Only the method positions on the rank axis are interpreted descriptively.

Compute resources. The experiments were executed primarily on a local workstation running Ubuntu 24.04 LTS, equipped with an AMD Ryzen 9 3950X CPU (16 cores and 32 threads), 128 GB of RAM, and one NVIDIA GeForce RTX 5090 GPU with 32 GB of VRAM. Exact wall-clock times and the additional compute consumed by preliminary experiments were not systematically recorded; runtimes may also vary with software configuration and available parallelism.

Reproducibility. An anonymized repository accompanying the submission contains the complete R–Python pipeline for dataset construction, fixed series-level splits, model fitting and evaluation, Chronos-2 inference, sensitivity analysis, and regeneration of the reported tables and figures. The scripts record the seeds, window settings, model hyperparameters, checkpoint identifiers, and adjustment grid.

The README documents installation and the exact execution order. Versioned requirements, Conda specifications, and Docker definitions specify the software environments, while generated artifacts use fixed project paths.

Asset licences. The accompanying code developed for this study is released under the MIT License. The M4 data are accessed through `M4comp2018` version 0.2.0 (GPL-3); the `amazon/chronos-2` and `paris-noah/Mantis-8M` checkpoints are released under Apache-2.0; `sktime` is released under BSD-3-Clause; `XGBoost` is released under Apache-2.0; and `scmamp` is released under GPL-2. Exact software versions and checkpoint identifiers are recorded in the accompanying repository. These assets are credited to their original authors, used under their respective terms, and downloaded from their official sources rather than redistributed.

Code availability. The anonymized repository is available to reviewers, and a cleaned public version will be released after the review period.

Broader Impacts. Directional forecasting may improve decision support in domains where anticipating the direction of change is more important than predicting its precise magnitude. However, incorrect directional predictions or deployment beyond the evaluated M4 setting could lead to harmful financial or operational decisions. Domain-specific validation, uncertainty assessment, monitoring, and human oversight are therefore necessary, particularly in high-stakes applications.

NeurIPS Paper Checklist

1. **Claims**

   Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

   Answer: [Yes]

   Justification: The abstract and Section 1 accurately describe the framework, descriptive M4 analysis, and post hoc proof-of-concept adjustment, as supported by Sections 3–6 and Appendices A–B.

   Guidelines:

   - The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
   - The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
   - The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
   - It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. **Limitations**

   Question: Does the paper discuss the limitations of the work performed by the authors?

   Answer: [Yes]

   Justification: Sections 3–6 and Appendices A, B, and D discuss the post hoc, non-held-out analysis, descriptive and setting-dependent results, pretraining overlap, and restriction to M4 and the evaluated methods.

   Guidelines:

   - The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
   - The authors are encouraged to create a separate "Limitations" section in their paper.
   - The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
   - The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
   - The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
   - The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
   - If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
   - While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. **Theory assumptions and proofs**

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: Section 3 and Appendix C present definitions and procedural equations, but the paper contains no theoretical results or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. **Experimental result reproducibility**

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Appendix D and the accompanying anonymized repository document the complete experimental pipeline, configurations, software environments, and procedures for reproducing the reported results.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general. releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in

some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. **Open access to data and code**

   Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

   Answer: [Yes]

   Justification: Appendix D and the accompanying anonymized repository provide access to the M4 data, environment specifications, exact execution commands, and scripts for reproducing the reported experiments, tables, and figures.

   Guidelines:

   - The answer [N/A] means that paper does not include experiments requiring code.
   - Please see the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
   - While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
   - The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (`https://neurips.cc/public/guides/CodeSubmissionPolicy`) for more details.
   - The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
   - The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
   - At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
   - Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. **Experimental setting/details**

   Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

   Answer: [Yes]

   Justification: Section 4, Appendix C, and Appendix D specify the data splits, training and evaluation procedures, model configurations, hyperparameter selection, and accompanying implementation details.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
   - The full details can be provided either with the code, in appendix, or as supplemental material.

7. **Experiment statistical significance**

   Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

   Answer: [No]

   Justification: Point-forecast accuracy follows the official M4 evaluation protocol, with sMAPE, MASE, and OWA computed over the complete official forecast horizon, as described in Section 4 and Appendix A. These fixed-evaluation-set results are treated as descriptive; no error bars, confidence intervals, or fully specified inferential tests are reported for the main claims.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. **Experiments compute resources**

   Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

   Answer: [Yes]

   Justification: Appendix D reports the Ubuntu 24.04 hardware used for the experiments, including the compute workers, memory, and storage. The complete reproducibility pipeline typically requires approximately one week of wall clock time on this system. Runtime varies because some R and Python workloads cannot always be executed reliably in parallel and may require sequential execution or restarting. The pipeline uses cached intermediate outputs and safe restart procedures, so completed work is preserved. This estimate excludes preliminary and failed development runs, which were not systematically recorded.

   Guidelines:

   - The answer [N/A] means that the paper does not include experiments.
   - The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
   - The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
   - The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. **Code of ethics**

   Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics `https://neurips.cc/public/EthicsGuidelines`?

   Answer: [Yes]

   Justification: The research conforms to the NeurIPS Code of Ethics. It uses publicly available benchmark data, involves no human participants or sensitive personal data, and follows appropriate standards of research integrity, transparency, and reproducibility.

   Guidelines:

   - The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.

- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

    Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

    Answer: [Yes]

    Justification: Appendix D discusses the potential decision-support benefits of directional forecasting, the harms arising from incorrect or inappropriate deployment, and corresponding mitigation measures.

    Guidelines:

    - The answer [N/A] means that there is no societal impact of the work performed.
    - If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
    - Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
    - The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
    - The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
    - If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. **Safeguards**

    Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

    Answer: [N/A]

    Justification: The work does not release a high-risk model, scraped dataset, or sensitive data. It uses the publicly available M4 benchmark and existing pretrained models; therefore, specialised release safeguards are not applicable.

    Guidelines:

    - The answer [N/A] means that the paper poses no such risks.
    - Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
    - Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
    - We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. **Licenses for existing assets**

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Appendix D identifies the versions and licences of the dataset, pretrained models, and principal software assets, which are appropriately credited and used under their respective terms.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The new experimental code and reproducibility pipeline are documented in Appendix D and the accompanying anonymized repository, which provides installation instructions, execution commands, environment specifications, asset licences, and descriptions of generated outputs. The work introduces no human-derived assets requiring consent.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The research uses the existing M4 benchmark and does not involve crowdsourcing, human participants, or human-subject data collection.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.

- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [N/A]

    Justification: The research does not involve crowdsourcing, human participants, or human-subject data; therefore, IRB or equivalent institutional approval was not required.

    Guidelines:

    - The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
    - Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
    - We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
    - For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

    Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

    Answer: [N/A]

    Justification: Neither LLMs nor agents were used as an important, original, or nonstandard component of the methodology.

    Guidelines:

    - The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
    - Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.